
Using CNNs and Hyperparameter-Optimized Boosting Techniques to Predict Vital Plant Traits

Prem Patel

School of Computer Science
University of Waterloo
Waterloo, ON, N2L 3G1
p352pate@uwaterloo.ca
August 9, 2024

Abstract

A pretrained Resnet50 convolutional model was used to extract 2048 key photographic features of plant photographs, captured through citizen science. Pairing this with 163 ancillary features allowed for a rich feature set, which was then pipelined into a tuned boosting algorithm (XGBoost) to accurately predict target features of unseen images. Ultimately, this resulted in an average R^2 test score of approximately 0.3 across the six target features, showcasing successful model performance. See GitHub for the final source code used to achieve this task, using libraries such as PyTorch for ResNet feature extraction, XGBoost for boosting, and scikit-optimize for Bayesian hyperparameter tuning.

1 Introduction

The primary task was to predict plant properties (traits) using plant photographs along with supporting numerical data. A large training database of 43,363 images was made available, along with 6,391 test images to evaluate model performance. The accuracy of these model predictions were evaluated using an R^2 score, with 6 target variables to be predicted. With access to this large dataset, it was necessary to employ a rich feature extraction method, i.e. ResNet50. To produce as robust a model as possible, it was also necessary to utilize the supporting numerical data. This started off as a combined neural network using both the ResNet model (with the final connected layer removed) and ancillary data as inputs. After thorough exploration of various architectures and a re-examination of the problem, it was clear that a boosting strategy was more desirable. This resulted in a final model that utilized XGBoost, an effective boosting algorithm to combine many weak learners to create a more robust model for trait prediction.



Figure 1: Example of plant photograph used for model training

22 2 Related Works

23 Given the photographic nature of the training data, it was natural to consider using a convolutional
 24 network in some capacity. This approach has been used in the past, where a fully connected ResNet
 25 CNN was used to predict plant traits from similar photographic data, with bioclimatic data being fed
 26 into the model as well (Schiller et al., 2021). This work was done by Christopher Schiller, Sebastian
 27 Schmidtlein, Coline Boonman, Alvaro Moreno-Martínez & Teja Kattenborn.

28 A key difference from this model is the removal of the final layer, as the ResNet model was simply
 29 used as a means of extracting features in this case, rather than generating final predictions. As such,
 30 ancillary data was concatenated and then fed to the XGBoost algorithm, generating predictions on
 31 a trait-by-trait basis. Another key difference is that hyperparameter tuning was utilized to optimize
 32 the performance of this boosting strategy.

33 The key limitation of the past approach is that it is comparatively less robust and heavily dependent
 34 on data preprocessing to achieve positive test results. With one network responsible for predicting
 35 6 differing traits, along with features with greatly varied scales, a boosting approach was deemed to
 36 be more appropriate as it would iterate through each target trait and use a decision-tree approach to
 37 decide which features are more relevant to a particular trait.

38 See below for the pseudocode outlining the XGBoost algorithm methodology in determining optimal
 39 decision tree splits (Chen & Guestrin, 2016).

Algorithm 1: Exact Greedy Algorithm for Split Finding

Input: I , instance set of current node

Input: d , feature dimension

```

1  $gain \leftarrow 0$ ;
2  $G \leftarrow \sum_{i \in I} g_i, \quad H \leftarrow \sum_{i \in I} h_i$ ;
3 for  $k = 1$  to  $m$  do
40 4  $G_L \leftarrow 0, \quad H_L \leftarrow 0$ ;
5   for  $j$  in  $sorted(I, \text{by } x_{jk})$  do
6      $G_L \leftarrow G_L + g_j, \quad H_L \leftarrow H_L + h_j$ ;
7      $G_R \leftarrow G - G_L, \quad H_R \leftarrow H - H_L$ ;
8      $score \leftarrow \max \left( score, \frac{G_L^2}{H_L + \lambda} + \frac{G_R^2}{H_R + \lambda} - \frac{G^2}{H + \lambda} \right)$ ;

```

Output: Split with max score

41 3 Main Results

42 See below for pseudocode detailing the entire learning pipeline, as implemented fully at this GitHub
43 link. Design decisions and relevant insights into model adaptation will then be explained.

Algorithm 2: Plant Trait Prediction Pipeline (Simplified)

Input: *train_data*, Numerical training data
Input: *test_data*, Numerical test data
Input: *images_train*, Training images
Input: *images_test*, Test images
Input: *tuning_flag*, Hyperparameter tuning method ('G', 'B', or 'N')
Output: *submission_file*, Predictions file

- 1 **1. Define Data Transformations ;**
- 2 **Function** GetTransforms () :
- 3 **2. Preprocess Data ;**
- 4 **Function** Preprocess (*train_data*, *test_data*) :
- 5 **3. Prepare DataLoaders ;**
- 6 **Function** CreateDataLoaders (*images_train*, *images_test*) :
- 7 **4. Extract Features ;**
- 8 **Function** ExtractFeatures (*train_loader*, *test_loader*) :
- 9 **5. Combine Features ;**
- 10 **Function** CombineFeatures () :
- 4411 **6. Set Up XGBoost and Tuning Models ;**
- 12 **Function** InitializeXGBoost () :
- 13 **7. Train and Tune XGBoost Model for Each Target Trait ;**
- 14 **Function** TrainAndTuneModels () :
- 15 **if** *tuning_method* = 'G' **then**
- 16 | **Function** GridSearch () :
- 17 **else if** *tuning_method* = 'B' **then**
- 18 | **Function** BayesianSearch () :
- 19 **else**
- 20 | **Function** XGBoostManual () :
- 21 **8. Generate Predictions ;**
- 22 **Function** Predict (*test_data*) :
- 23 **9. Save Results ;**
- 24 **Function** SavePredictions (*submission_file*) :
- 25 **10. Notify Results ;**
- 26 **Function** SendEmail (*submission_file*) :

45 3.1 Problem Formulation

46 The problem of predicting plant traits from photographs was framed as a supervised regression prob-
47 lem. The challenge was to map complex visual features from plant images to numerical target traits,
48 leveraging both image-based features and ancillary data. Mathematically, this can be represented
49 as learning a function $f : X \rightarrow Y$, where X is the combined feature space of image features and
50 ancillary data, and Y is the set of target traits. The goal was to find a function that minimizes the
51 mean squared error between predicted and actual trait values across six target traits.

52 3.2 Preliminary Models and Preprocessing

53 Initial experiments involved designing a fully connected neural network that integrated features from
54 a convolutional neural network (CNN) and ancillary data. Ancillary preprocessing steps included
55 log transformation of target values to stabilize variance, followed by min-max normalization of
56 both training and test datasets, to ensure uniform scaling. Outlier removal consisted of eliminating
57 data points exceeding three standard deviations from the mean, to mitigate their impact on model
58 performance. Images were also preprocessed through 224x224 pixel resizing, normalization and
59 tensor conversion. Training images were also augmented at random, using flipping and colour jitter.
60 Despite these efforts, preliminary models were prone to overfitting and achieved an average R^2
61 score below 0.2 on unseen test images. Regularization techniques including early stopping, batch
62 normalization and dropout were used but did not significantly improve model generalization.

63 3.3 Boosting and Hyperparameter Tuning

64 A major improvement came with the incorporation of XGBoost, a boosting algorithm known for
65 its robustness and performance. Preprocessing and augmentation remained the same as above,
66 now with the CNN and ancillary data being used strictly as feature inputs to the boosting process.
67 Bayesian hyperparameter tuning was used to optimize the model, exploring a broad search space
68 that balanced exploration with practical computational constraints. This approach led to a signifi-
69 cant improvement, with the final model achieving an average R^2 score of 0.3. Hyperparameter tuning
70 involved fine-tuning parameters like learning rate, number of estimators, and tree depth, which were
71 crucial for improving model accuracy.

72 3.4 Ablation Studies

73 Various configurations were tested to assess their impact on model performance. Different pre-
74 processing methods (e.g., log-transforming targets vs. all training data), normalization techniques
75 (z-score vs. min-max), and architecture choices (e.g., ResNet18 vs. ResNet50) were compared. The
76 effectiveness of combining CNN features with ancillary data versus using them separately was also
77 evaluated. The switch to a CNN-XGBoost pipeline was motivated by the need for a more robust
78 solution, where tuning could be used to optimize model performance. Various tuning methods were
79 implemented (Bayesian search, Grid search, no tuning) for further ablation.

80 4 Conclusion

81 Experiments indicated that a hybrid approach combining CNN and ancillary feature extraction with
82 tuned XGBoost for regression can significantly enhance prediction accuracy for plant traits. The
83 key takeaway is the importance of choosing the right architecture and preprocessing strategy suited
84 to the data. Observations revealed that preprocessing choices, such as log transformation and nor-
85 malization, significantly influenced performance of the preliminary, fully connected model. Many
86 of these choices had a negligible impact for boosting. In retrospect, starting with a robust boosting
87 approach from the beginning would have allowed for deeper exploration of model variations, while
88 reducing the need for excess testing of preprocessing techniques.

89 4.1 Future Directions

90 Constraints of time and computational resources were notable, impacting the breadth of exploration.
91 With additional resources, expanding the search space for hyperparameter tuning could yield even
92 better results. Exploring more sophisticated feature extraction techniques might further enhance
93 model performance, such as applying principal component analysis to reduce dimensionality. Bal-
94 ancing computational feasibility with model complexity remains a crucial consideration in machine
95 learning projects.

96 **Acknowledgement**

97 Thank you to Christopher Schiller, Sebastian Schmidlein, Coline Boonman, Alvaro Moreno-
98 Martínez & Teja Kattenborn for providing initial inspiration, as well as Yao-Liang Yu for providing
99 the academic inspiration necessary in producing this work.

100 **References**

- 101 Chen, T. and C. Guestrin (2016). “XGBoost: A Scalable Tree Boosting System”. In: *Proceedings*
102 *of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*.
103 KDD ’16. ACM.
- 104 PyTorch (2024). “resnet50: Models and pre-trained weights”. Accessed: 2024-08-01.
- 105 Schiller, C., S. Schmidlein, C. Boonman, A. Moreno-Martínez, and T. Kattenborn (2021). “Deep
106 learning and citizen science enable automated plant trait predictions from photographs”. *Scientific*
107 *Reports*, vol. 11.
- 108 XGBoost Contributors (2018). “XGBoost Parameters”. Accessed: 2024-08-05.
- 109 Yu, Y. (2024). “CS480/680: Introduction to Machine Learning”. Accessed: 2024-07-26.